

How Long Is a Thought? The Stopwatch and the Thinker Are Measuring Different Things

Human-readable companion to How Long Is a Thought? Separating Observer-Measured Duration from Temporal Use and Reconstruction

Human-readable companion to the completed scientific manuscript. This version preserves the scientific claims and limits of the completed manuscript while presenting the motivating problem and argument in more direct language. It does not constitute a separate research contribution.

Abstract

A thought can be timed from the outside. If an experimenter defines when a cognitive episode begins and ends, a stopwatch gives that bounded episode an elapsed duration. But the stopwatch does not reveal whether metric elapsed time was represented or causally used by the person while the thought was occurring. A later answer to “how long did that take?” does not settle the question either, because the answer may be reconstructed after the query from remembered events, bodily or environmental cues, priors, retained temporal information, or mixed routes. This paper separates observer-measured duration, latent online temporal use, post-query temporal basis or reconstruction, and the reports produced after each stage. It proposes a minimal protocol: preserve one temporally naïve thought, obtain one quick retrospective estimate, deliberately ask the participant to think about the duration, time that reflection episode, then ask for a quick estimate of the reflection episode itself. Crucially, the protocol does not ask for a second estimate of the original thought. The recursive step demonstrates that the same objective-duration/report identification problem reappears one level higher. The framework does not identify a timing mechanism; it shows which measurements must not be collapsed.

1. The question sounds simple because it hides several questions

A person starts thinking and later signals that the thought is complete. From the experimenter’s side, the interval can be timed. If the start and stop markers are declared in advance, the result is an ordinary elapsed duration for the bounded episode.

But “how long did the thought take?” can also ask something about the thinker. Did metric elapsed time participate in the thought while it was happening? Was temporal information available but irrelevant to the computation? Did the duration become inferable only after someone asked? Was the final number reconstructed from the amount that happened rather than read from an online clock-like representation?

Those are different scientific objects. The paper’s central discipline is to keep them different even when they happen to agree numerically.

2. What the stopwatch establishes—and what it cannot

The observer-measured duration, TO, is the elapsed time between the experiment’s declared onset and offset. It tells us how long the operationally bounded episode lasted. It does not tell us what temporal information the person represented, accessed, or used during that interval.

This is not a claim that the thought somehow occurred outside physical time. Every cognitive process unfolds in time. The question is narrower: did metric elapsed-time information causally participate in the ongoing thought? The manuscript calls that latent relation UT.

$$\text{observer duration TO} \neq \text{latent online temporal use UT}$$

The distinction also does not treat a key press as the metaphysical beginning or end of a natural thought. “A thought” is operationalized as a bounded cognitive episode. Natural cognition can contain inner speech, imagery, memories, subgoals, and nested transitions. The experiment must declare its boundary and limit the claim to the episode so defined.

3. A later estimate creates a new measurement stage

Suppose the stopwatch reads 11.8 seconds and the participant says “about 12 seconds.” That close match is informative: the post-query judgment contains useful information about the observer-defined interval. It still does not tell us how the estimate was produced.

Timing research already distinguishes prospective timing, where people know in advance that duration matters, from retrospective timing, where the duration query arrives unexpectedly. Retrospective estimates can depend on memory and event structure, while prospective instructions can recruit attention toward time and thereby change the process being studied (Block & Zakay, 1997; Grondin, 2010; Yue et al., 2026).

The first post-query estimate is therefore not treated as a passive readout. Asking the question is an intervention on subsequent cognition. The participant can search memory, use bodily or environmental cues, map a vague sense of extent onto clock units, consult retained temporal information, or combine several routes.

An inaccurate answer is equally non-identifying. It could reflect absent online temporal use, but it could also reflect later reporting noise, scale-mapping error, uncertain boundaries, bias, or poor access to information that did exist online. “I have no idea” establishes non-reportability under that probe, not temporal absence.

4. The six objects in the minimal architecture

Object	Meaning	Observed directly?
TO	Observer-measured duration of the original bounded episode	Yes
UT	Whether metric elapsed time causally participates during the naïve episode	No; central latent target
Apost	Post-episode temporal basis/access/reconstruction available after the query	Model-dependent category
$\hat{T}Q$	First quick duration estimate of Episode 1 after the surprise query	Yes
TR	Observer-measured duration of the deliberate reflection episode	Yes
$\hat{T}R$	Quick estimate of the duration of the reflection episode	Yes

The framework does not define Apost as the opposite of UT. Temporal information could be represented online but unused, used online but later difficult to report, accessed only after the query, reconstructed from non-temporal cues, or combined across routes. The four observables therefore constrain models without uniquely identifying the latent organization.

5. The minimal naïve-thought protocol

The protocol is designed to preserve one episode in which the participant has not yet been told that elapsed duration is the target. The first task should be bounded and non-temporal: it should not ask the participant to count, estimate duration, wait a set interval, synchronize to a rhythm, or stop according to time.

Step	What happens	What is recorded
1	Episode 1 occurs while the participant is temporally naïve to the study target.	TO
2	After Episode 1: “Give me a quick estimate of how long that took.”	$\hat{T}Q$, with response latency optionally recorded
3	Then: “Now think about how long it took.” A new timer starts.	Episode 2 begins; time is now explicitly relevant
4	Participant indicates when the deliberate reflection is complete.	TR

Step	What happens	What is recorded
5	Participant quickly estimates how long Episode 2 itself took.	$\hat{T}R$; sequence stops

The word “quick” is intentionally modest. It means less elaborated than the deliberate reflection that follows; it does not mean reconstruction-free, instantaneous, or guaranteed to expose a pre-existing duration representation.

The first surprise query is a scarce resource. Once duration has been asked about, later trials occur in a different informational regime. Repeated trials can still be useful, but they should be treated as temporally informed or prospective rather than pretending the original naïveté can be restored.

6. Why the protocol deliberately does not ask for a second estimate of Episode 1

After $\hat{T}Q$, the participant already knows the temporal target and has committed to a first answer. The instruction to think about the duration then creates an explicitly temporal reflection episode. A second estimate of Episode 1 would be vulnerable to anchoring on the first answer, consistency pressure, demand effects, and deliberate reconstruction.

More importantly, that second estimate would not repair the original identification problem. A change after reflection would show that intervening cognition changed the report, but it would not by itself reveal whether metric elapsed time was causally used during the original naïve thought.

The final protocol therefore does something cleaner: it makes Episode 2 the new target. The experimenter times the reflection itself and then asks how long that reflection took.

7. The recursive probe repeats the same identification problem

Episode 2 is unusual because temporal duration is now explicitly the topic of thought. The participant may replay events, count remembered steps, consult bodily cues, compare against familiar durations, retrieve temporal information, or use a mixture of strategies. TR tells us only how long that reflection episode lasted between its markers.

When the participant then estimates Episode 2 as $\hat{T}R$, agreement with TR still does not prove that metric elapsed time was causally monitored throughout Episode 2. Disagreement does not prove temporal information was absent. The same observer-duration/report non-identification has reappeared at the next level.

That recurrence is the point of the recursive probe. It makes the inferential structure visible; it is not a unique mechanistic discriminator. One recursive level is enough. Repeating the procedure indefinitely would turn the task into expected timing, not reveal a metaphysical regress.

8. What ordinary report patterns can and cannot tell us

Observation	Warranted inference	Inference not warranted by that observation alone
$\hat{T}Q$ closely matches TO	The first post-query report contains information about the observer-defined interval.	UT was causally used online.
$\hat{T}Q$ differs greatly from TO	Objective and reported duration differ.	No temporal information existed or was used online.
Participant says “I have no idea.”	Duration was not reportable under this probe.	Temporal information was absent, inaccessible online, or unused.
$\hat{T}R$ closely matches TR	The report about Episode 2 tracks its observer-defined duration.	Episode 2 continuously monitored metric time.
Reflection adds little interpretable structure	The staged protocol may add little beyond an ordinary retrospective query.	TO and later report become identical measurement objects.

The cleanest evidence for UT would require an additional independently validated online manipulation or measure that is sensitive to metric elapsed-time information during Episode 1, preserves the relevant naïveté, and

makes competing accounts predict different outcomes. The minimal protocol deliberately does not pretend to supply that missing bridge.

9. Edge cases that change the question

If the thought itself is “count ten seconds,” “decide when a minute has passed,” or “stop after roughly 15 seconds,” time is relevant by construction. Such tasks can validate timing competence but are not clean tests of naïve non-temporal use.

External clocks, periodic sounds, screen animation, breathing cycles, heartbeat, and salient event sequences can provide temporal cues. Eyes-closed performance is one optional way to reduce some visual cues, not part of the theory. A good study manipulates or measures relevant cues rather than assuming they are absent.

Very short episodes make onset/offset and motor-latency errors proportionally large. Very long episodes accumulate memory, context change, boredom, and event boundaries that support reconstruction. The initial duration range should therefore be empirically calibrated rather than dictated by the framework.

Spontaneous time monitoring is also possible. A participant may attend to time without being instructed. A retrospective question about whether this happened can be informative but is still retrospective; stronger work would need converging measures.

10. What existing work already owns

The components are not presented as new discoveries. Prospective versus retrospective timing is established (Block & Zakay, 1997; Grondin, 2010). Memory and event structure can shape retrospective duration (Yue et al., 2026). Quantified introspection can track objective processing while missing specific delays (Corallo et al., 2008; Marti et al., 2010). First-order temporal judgments and second-order temporal metacognition are already active empirical topics (Cropper et al., 2024; Pavailler et al., 2025). Experience-sampling work already deals with reactivity and difficult thought boundaries (Hurlburt & Akhter, 2006; Klinger, 2013).

The residual contribution is the integrated object-typing and staged sequence: one naïve bounded episode, one minimally elaborated post-query estimate, one explicitly temporal reflection episode, and one estimate of that new episode—without treating any report as automatic proof of online temporal use.

11. How the protocol can fail without erasing the identification result

The constructive protocol can prove empirically unhelpful. Episode 2 may add no stable or interpretable information beyond the first retrospective estimate. TR and $\hat{TR}$ may be too noisy for useful analysis. Operational boundaries may be poor. The queries may perturb cognition so strongly that the staged measures are dominated by reactivity.

Those outcomes would narrow or eliminate the value of this specific protocol. They would not make the stopwatch duration and the later report identical to latent online temporal use. The deeper non-identification claim is challenged only by an independently justified bridge or assumptions that legitimately collapse the latent alternatives.

12. Conclusion

For an observer, a bounded thought has a duration between declared markers. For the person having the thought, that number does not tell us whether metric time was represented, accessible, causally used, or only reconstructed after someone asked.

The minimal protocol keeps each transition visible. It times the naïve thought without asking about time, asks once for a quick estimate, deliberately creates a new duration-focused reflection episode, times that episode, and then asks for an estimate of the reflection itself. It stops there instead of contaminating the original target with a second estimate.

The answer to “how long is a thought?” therefore depends on which measurement object is being requested. The stopwatch can answer one version directly. A report answers another. Neither, by itself, identifies the causal role of time inside the thought.

References

- Block, R. A., & Zakay, D. (1997). Prospective and retrospective duration judgments: A meta-analytic review. *Psychonomic Bulletin & Review*, 4, 184–197. DOI: 10.3758/BF03209393.
- Corallo, G., Sackur, J., Dehaene, S., & Sigman, M. (2008). Limits on introspection: Distorted subjective time during the dual-task bottleneck. *Psychological Science*, 19, 1110–1117. DOI: 10.1111/j.1467-9280.2008.02211.x.
- Cropper, S. J., Little, D. R., Xu, L., Bruno, A. M., & Johnston, A. (2024). Measuring the perception and metacognition of time. *Journal of Vision*, 24(3), 5, 1–28. DOI: 10.1167/jov.24.3.5.
- Grondin, S. (2010). Timing and time perception: A review of recent behavioral and neuroscience findings and theoretical directions. *Attention, Perception, & Psychophysics*, 72, 561–582. DOI: 10.3758/APP.72.3.561.
- Hurlburt, R. T., & Akhter, S. A. (2006). The Descriptive Experience Sampling method. *Phenomenology and the Cognitive Sciences*, 5, 271–301. DOI: 10.1007/s11097-006-9024-0.
- Klinger, E. (2013). Goal commitments and the content of thoughts and dreams: Basic principles. *Frontiers in Psychology*, 4, 415. DOI: 10.3389/fpsyg.2013.00415.
- Marti, S., Sackur, J., Sigman, M., & Dehaene, S. (2010). Mapping introspection's blind spot: Reconstruction of dual-task phenomenology using quantified introspection. *Cognition*, 115, 303–313. DOI: 10.1016/j.cognition.2010.01.003.
- Nicolai, C., Chaumon, M., & van Wassenhove, V. (2024). Cognitive effects on experienced duration and speed of time, prospectively, retrospectively, in and out of lockdown. *Scientific Reports*, 14, 2006. DOI: 10.1038/s41598-023-50752-7.
- Pavailler, N., Gevers, W., & Burle, B. (2025). Temporal metacognition: Direct readout or mental construct? The case of introspective reaction time. *Journal of Experimental Psychology: General*, 154, 1122–1148. DOI: 10.1037/xge0001708.
- Phillips, I. (2012). Attention to the passage of time. *Philosophical Perspectives*, 26, 277–308. DOI: 10.1111/phpe.12007.
- Phillips, I. (2013). Perceiving the passing of time. *Proceedings of the Aristotelian Society*, 113, 225–252. DOI: 10.1111/j.1467-9264.2013.00353.x.
- Sutton, J. (2011). Time, experience, and descriptive experience sampling. *Journal of Consciousness Studies*, 18(1), 118–129.
- Vitulli, W. F., & Shepard, H. A. (1996). Time estimation: Effects of cognitive task, presentation rate, and delay. *Perceptual and Motor Skills*, 83, 1387–1394. DOI: 10.2466/pms.1996.83.3f.1387.
- Yue, W. W. Y., Liu, J., Yao, Z., Zhang, Y., Mu, Z., & Hu, X. (2026). Retrospective duration judgments of naturalistic events depend on memories of event boundaries. *Psychonomic Bulletin & Review*, 33, 33. DOI: 10.3758/s13423-025-02833-z.