

Weight: Why the Same Words Do Not Carry the Same Consequence

Abstract

People often speak as though communicating a statement also communicates what that statement means for them. Yet two people can possess the same content, understand the same proposition, and endorse it as true while it carries different consequences for attention, belief revision, memory, valuation, policy, or action. This theoretical article proposes Weight as a disciplined, human-readable identification language for such consequentiality, not as a universal mechanism or common psychological unit. A scientific Weight claim must identify the already-defined object, target system, local consequence relation, context, and evidence; descriptive, predictive, and causal claims require correspondingly stronger evidence. The framework separates content possession, generative understanding, local consequence, provenance, observer estimate, and communication transfer, and forbids automatic inference among them or across unlike numerical scales. Established work on communication, relevance, involvement, attitude strength, shared reality, epistemic vigilance, source monitoring, risk, empathy, and measurement provides the component machinery. Weight survives scientifically only if its typed declaration standard improves measurement or inference beyond a matched generic-explicitness control; otherwise it should contract to useful ordinary language.

Keywords: consequentiality; communication; understanding; provenance; belief updating; shared reality; measurement

1. THE ORDINARY PROBLEM

"I told you."

That sentence often means more than "the words left my mouth." It often means: I communicated the seriousness of this. You knew what this meant. You should have understood the consequences. You should therefore have responded as I would have responded.

Those are different claims.

A person can successfully transmit a proposition without transmitting what that proposition carries for them. A listener can repeat the statement accurately, explain its literal meaning, and even agree that it is true while the statement remains comparatively light in the listener's own system. The sender may then experience the listener's response as indifference, denial, stupidity, dishonesty, or refusal. Sometimes those descriptions are correct. Sometimes the more basic problem is that the same content was present on both sides while its consequences were not.

The reverse problem is equally familiar. Someone can hear a statement that is true and understand the sentence without receiving it as an immediate fact. A stranger says, "Your building is on fire." A physician says, "If you keep doing this, you may die." A partner says, "I am leaving." A manager says, "Your position is being eliminated." A family member says, "He is gone."

The words may be clear. Yet the receiving system may initially treat them as remote, provisional, unreal, abstract, exaggerated, low-confidence, or simply not yet integrated with everything else it implies. People describe this afterward with phrases such as "I knew, but it had not hit me," "I heard them, but I did not understand how serious it was," or "I understood what they were saying, but it did not feel real yet."

These phrases are scientifically dangerous if interpreted literally. Feeling is not truth. A proposition does not become factual because it produces emotion. Nor does a proposition

become false because it initially produces little emotion. But the phrases point to a legitimate distinction: a representation can be present before its consequences are organized strongly enough to alter attention, expectation, memory, valuation, policy, or action.

Everyday language often calls that difference weight.

The purpose of this paper is to make that word usable without allowing it to become an explanation for everything.

2. WHAT WEIGHT MEANS

The proposed human-language definition is:

WEIGHT is the consequence something carries for a system, relative to what that thing is, the relation being asked about, and the situation in which it is encountered.

Several parts of this definition matter.

First, Weight belongs to something. A belief can have Weight. A warning can have Weight. A memory can have Weight. An observation can have Weight. An intention can have Weight. An anticipated consequence can have Weight. The word does not supply the identity of the thing being discussed. In this paper, calling the object "typed" means only that the object is named using the science that already defines it rather than being redefined by Weight.

Second, Weight is system-relative. The same statement may carry different consequence for the speaker and listener, for two listeners, or for the same person at different times.

Third, Weight is relation-relative. Saying that something "has Weight" is scientifically incomplete unless one can say what kind of consequence is being claimed. Does it bear on attention? Memory priority? Credence? Emotional response? Choice policy? Persistence? Avoidance? Physiological response? Social trust? A mature local science may already have a better name and measure for that particular relation. Weight does not replace it.

Fourth, Weight is context-relative. A memory that is nearly irrelevant in one setting can become highly consequential in another. A risk that feels remote in the morning can become behaviorally decisive after new evidence arrives. A belief can retain the same verbal content while its practical consequence changes.

Fifth, consequence in the human definition is not itself a causal claim. Scientific Weight statements must match their evidential level. A descriptive Weight claim may be supported by a validly measured association or state description. A predictive Weight claim requires demonstrated predictive performance under a declared target and generalization regime—for example, prospective prediction, held-out evaluation, or cross-validation appropriate to the design. A causal Weight claim requires an intervention or another valid causal-identification strategy. Weight must not turn an observed endpoint into retrospective proof of a causal route.

This makes Weight an interface term rather than a universal substance. It is useful precisely because people already use the word across domains. The scientific requirement is that each use be attached back to an already-defined object and a declared consequence relation.

Weight therefore does not mean "importance" in the abstract. Importance is one established construct that may be the correct local meaning in some cases. The attitude-strength literature, for example, shows that personal importance and related strength properties have measurable relationships with information processing, memory organization, emotional response, behavioral intention, and action while also showing that attitude strength is not a single simple property (Krosnick et al., 1993; Howe & Krosnick, 2017). Weight should inherit such local distinctions rather than erase them.

The interface earns scientific use only when making the declaration changes something: what is measured, what can be inferred, what can be compared, what must be controlled, or what

would count as a failed prediction. If it changes none of those, "Weight" remains ordinary language rather than a scientific contribution.

3. WHAT WEIGHT DOES NOT MEAN

The usefulness of Weight depends more on what it forbids than on what it names.

Weight is not truth.

A false belief can carry enormous Weight. A true statement can carry very little. The truth status of a proposition and its consequence for a system are different questions.

Weight is not belief.

A person can assign high probability to a proposition while treating it as practically remote. A person can also organize behavior around a possibility they do not consider especially likely when the cost of being wrong is high. Belief and consequence can interact without being identical.

Weight is not understanding.

Someone can understand a mechanism deeply without caring about it. Someone can care intensely about something they understand poorly. Research on explanatory understanding demonstrates that people's sense of understanding can diverge substantially from their ability to generate an explanation (Rozenblit & Keil, 2002). Understanding therefore requires its own measurement.

Weight is not emotion.

Emotion may itself be the typed object, one consequence relation, or one contributor to another. But a statement can carry substantial decision consequence with little overt feeling, and strong feeling does not establish decision authority.

Weight is not authority.

Something may be highly consequential and still lose to competing goals, beliefs, constraints, incentives, habits, or policies. High Weight in one declared relation does not imply that the object will determine the selected action.

Weight is not salience.

What captures attention is not necessarily what controls judgment or action, and what controls action need not remain phenomenally salient.

Weight is not provenance.

Where a consequence came from is different from the consequence it currently carries. Direct experience, testimony, training, analogy, memory, social learning, and current bodily state can all contribute to current organization. They do not guarantee a particular present Weight.

Weight is not a causal effect by definition.

A relation can be measured without its causal direction being identified. Causal Weight claims require causal evidence rather than the ordinary-language intuition that something "mattered."

Weight is not behavior.

An endpoint cannot by itself tell us which represented objects were present, which participated, which lost competition, or which were irrelevant. Weight must not be defined retrospectively as "whatever caused the behavior."

These exclusions are not side notes. They are the reason to use a typed Weight framework at all.

4. THE FIRST COMMUNICATION FAILURE: CONTENT DOES NOT CARRY ITS WEIGHT WITH IT

Suppose a person says, "I know this will hurt you."

The sentence can be generated from many different organizations. The speaker may know the dictionary-level meaning of the relevant harm. They may possess a sophisticated causal model of what will happen. They may have personally experienced something similar. They may currently feel strong empathic distress. They may have no comparable experience but understand the other person's situation accurately. They may recognize the likely pain and continue anyway because another goal carries greater decision consequence. They may also be wrong about what the other person will experience.

The sentence alone does not identify which of those organizations produced it.

The same is true for the receiver. Hearing the sentence does not reproduce the speaker's internal organization. Language can transfer content with remarkable efficiency, but it does not copy a nervous system, autobiographical history, current body state, valuation structure, or policy hierarchy.

This yields the paper's first central claim:

**COMMUNICATION OF CONTENT DOES NOT ESTABLISH
COMMUNICATION OF WEIGHT.**

That claim does not require a mysterious transfer mechanism. It requires only that the things being transferred be measured separately.

A listener may acquire the proposition but not the model behind it. A listener may acquire the model but not the same consequence. A listener may infer that the topic carries great Weight for the speaker without acquiring that Weight personally. Or the listener's Weight may change dramatically for a different reason than the speaker's.

Research on shared reality is directly relevant. Shared reality concerns experienced commonality with another person's inner state about a target referent and can support

interpersonal connection and epistemic functions (Echterhoff, Higgins, & Levine, 2009). That literature demonstrates that perceived commonality itself matters. It does not require that two systems be internally identical. The distinction is important: a person can feel that another "gets it" because commonality has been established, while the exact sources and consequences on each side remain different.

The same caution applies to lived experience. Similar experience can alter empathic concern and perceived understanding without guaranteeing more accurate access to another person's individual state. Hodges and colleagues (2010) found that women with similar childbirth-related experience reported greater understanding and empathic concern, while effects on empathic accuracy were narrower. Israelashvili, Sauter, and Fischer (2020) reported that similar negative experience could even reduce emotion-recognition accuracy under some conditions. Lived experience can therefore provide a different reference structure without functioning as a certificate of understanding.

The practical implication is simple but consequential: "I told them" cannot establish "they received what this carries for me."

5. THE SECOND COMMUNICATION FAILURE: A STATEMENT CAN ARRIVE WITHOUT BECOMING IMMEDIATE REALITY

People often say that a statement "did not land" or "did not feel real." This paper does not treat immediate reality as a new psychological construct. The phrase compresses at least two scientifically different cases.

5.1 The proposition is not yet accepted

In the first case, the receiver does not yet accept the proposition as sufficiently credible. The relevant Weight belongs to evidence, testimony, or source credibility in relation to belief revision.

Imagine hearing "the building is on fire" from an unknown person walking past you on the street. Now imagine hearing the same sentence from a firefighter while a smoke alarm sounds and smoke is visible behind you. The proposition can be identical. Its truth does not change with the speaker. But the evidence available to the receiver changes dramatically, and therefore the proposition may receive very different immediate treatment.

Humans cannot rationally accept every communicated statement at full force. Communication exposes a receiver to error, misunderstanding, manipulation, and deception. Work on epistemic vigilance explicitly treats recipients as evaluating both communicated content and its source rather than simply absorbing testimony (Sperber et al., 2010). Communication-persuasion models similarly distinguish comprehension from acceptance and later action (McGuire, 1981). In Weight language, this is not generic "statement Weight." The parent object is Evidence, Testimony, Source, or another established epistemic object, and the declared relation is belief revision, confidence, or policy under uncertainty.

5.2 The proposition is accepted but weakly integrated elsewhere

The second case occurs after the receiver can already state, and perhaps sincerely endorse, the proposition. The problem is consequential uptake in some other relation.

A patient can believe that smoking increases mortality risk while the fact remains weakly connected to today's choice. A person can know that a relationship is ending while continuing to

plan as though the shared future remains intact. Someone can know that a parent is dying and still discover later that the fact has reorganized memory, priorities, time, and identity in ways that were not present when the sentence was first understood.

This is the structure behind "I knew, but it had not hit me yet."

The framework does not require us to claim that the earlier belief was fake. It permits a more precise statement:

**THE PROPOSITION WAS PRESENT BEFORE SOME OF ITS RELEVANT
CONSEQUENCES WERE PRESENT OR DETECTABLE.**

That wording is intentionally weaker than "all consequences." No observer has access to every consequence of a proposition, and many may never emerge. The scientific requirement is to name the relation under study.

Belief confidence, explanatory understanding, affective response, memory priority, valuation, and action policy are different measurable relations. A true proposition may alter one before another. Risk-as-feelings work provides a domain-specific example in which cognitive assessment and affective response can diverge (Loewenstein et al., 2001). The social-amplification-of-risk framework likewise shows that risk information can be processed and amplified or attenuated through communication and response systems (Kasperson et al., 1988). These are precedents for local dissociation, not a generic Weight mechanism.

The same caution applies to later reports of feeling. Two people can encounter the same consequential event and later show different conscious emotional trajectories, but a later report does not identify the upstream route. Weak or absent reported feeling is compatible with several organizations: relatively weak initial affective consequence; earlier transformation or implicit regulation; rapid integration leaving little sustained phenomenal residue; unresolved affective

information that becomes interpretable later; or phenomenal emotion that remains present but carries little authority for current reasoning or action. These are examples of observational equivalence, not mechanisms attributed by Weight. A downstream report cannot by itself establish when or by what route the consequence was generated, transformed, integrated, attenuated, or lost (Boden & Thompson, 2017; Gross, 1998; Gyurak et al., 2011).

The reward literature provides another useful precedent for this style of decomposition: liking, wanting, and learning are related but separable processes rather than one undifferentiated reward variable (Berridge, Robinson, & Aldridge, 2009). The point here is not that Weight is reward. The point is that mature science often advances by refusing to treat everyday unity as mechanistic unity.

Weight applies the same discipline to consequential uptake while keeping acceptance and downstream consequence separate.

6. HAVING THE WORDS, HAVING THE MODEL, HAVING THE CONSEQUENCE

The human-language center of the framework can be stated in one sentence:

**HAVING THE WORDS, HAVING THE MODEL, HAVING THE
CONSEQUENCE, AND TRANSMITTING ANY OF THEM ARE DIFFERENT
THINGS.**

For implementation, these can be kept separate without pretending they are new mechanisms.

6.1 Having the words

A person may be able to recognize, define, label, recall, or repeat the relevant proposition. Existing semantic and declarative-memory methods already measure such abilities. Nothing new is required.

Call this content possession when bookkeeping is needed.

6.2 Having the model

A person may be able to use the proposition relationally: explain why it is true, predict consequences, answer novel questions, identify boundary conditions, generate counterfactuals, or transfer the idea to a new situation. This is closer to what people commonly mean by understanding.

Rozenblit and Keil's (2002) work on the illusion of explanatory depth demonstrates why self-reported understanding cannot simply substitute for generative ability. Understanding should therefore be tested with tasks that require reconstruction, explanation, prediction, or transfer.

6.3 Having the consequence

The proposition or other object may actually bear on a defined local process: attention, memory, credence, valuation, emotional organization, policy, physiological state, persistence, avoidance, trust, or another measurable relation.

This is where Weight lives.

The important restriction is that the consequence relation must be named. "It had a lot of Weight" is a human statement. A scientific statement must add: Weight for what?

6.4 Transmitting any of them

Communication can change the receiver at several levels. The receiver may learn the words without the model, acquire a model without sharing the sender's consequence, or experience a consequence without fully reconstructing the sender's model. Communication can also change the receiver's estimate of the sender without changing the receiver's own Weight at all.

"Transfer" in this paper therefore does not mean that a quantity literally moves between two people. It means that some measurable property of the recipient changes after communication.

Communication theory already supplies important precedents for this separation. Speech-act theory distinguishes propositional content from what an utterance does, while grounding and common-ground work treats successful communication as coordinated evidence of mutual understanding rather than identity of internal states (Austin, 1962; Clark & Brennan, 1991). Persuasion models likewise separate exposure, comprehension, acceptance, retention, and action (McGuire, 1981). Weight does not claim those separations as new. It asks a different bookkeeping question across domains: after content is transmitted, which recipient-local consequence relation, if any, changed?

7. PROVENANCE: WHERE WEIGHT CAME FROM IS NOT THE SAME AS WHAT IT IS NOW

People care about where another person's understanding comes from.

"Have you been through this?"

"Do you know this because you studied it, or because it happened to you?"

"Are you telling me what you were taught, or what you saw?"

Those questions are not irrational. Source and history can change the structure available to a system. But provenance is easy to overclaim.

Source-monitoring research has long treated memory-source judgments as inferential and fallible rather than as direct readouts of origin (Johnson, Hashtroudi, & Lindsay, 1993). A person can remember content while misattributing its source. Conversely, different acquisition histories can produce similar present reports.

Weight therefore keeps provenance separate.

A direct experience can create durable reference structure, but direct experience does not guarantee current Weight. It does not guarantee correct interpretation. It does not guarantee empathy. It does not guarantee behavioral authority. It does not guarantee that another person with a similar experience will feel understood.

Provenance survives as scientifically important only when it changes what can be predicted or manipulated after present state has been adequately described. If two different histories produce the same present organization for the relation being studied, then the histories may remain relevant biography but need not remain separate causal variables for that question.

This rule preserves the intuitive importance of lived experience without turning lived experience into epistemic privilege.

8. TARGET-SIDE WEIGHT AND PERCEIVED WEIGHT

Communication introduces another distinction: what something carries for a target person or system and what another person thinks it carries for that target.

These are not guaranteed to match.

A listener may infer Weight from word choice, fluency, timing, prosody, respiration, facial behavior, narrative structure, disclosure, confidence, or context. Some of those cues may

be useful for estimating current internal states. None should be treated as a universal dictionary of provenance or truth.

The same outward pattern can arise from different sources: current arousal, personality, practiced communication, social skill, deliberate presentation, cultural norms, fatigue, anxiety, or experiential history. Therefore, perceived Weight is an estimate, not a measurement of hidden history.

The phrase target-side Weight does not imply privileged access to a true latent quantity. Empirical work only provides evidence about the declared target relation. A self-report, behavioral index, physiological measure, observer judgment, or model estimate is an observation with an evidential status, not automatic access to "actual Weight."

A researcher can therefore separately record:

- the target person's self-report relevant to the declared relation;
- behavioral, physiological, or other indices relevant to that relation;
- independently verified aspects of provenance when available;
- the observer's estimate of consequence or provenance;
- and the observer's confidence in that estimate.

The gap between those quantities may itself explain important social outcomes. Feeling understood may depend partly on what a person believes another person shares, not merely on what the other person actually shares. Shared-reality work makes this possibility important (Echterhoff et al., 2009), and research on felt understanding likewise distinguishes feeling understood from objective or externally assessed understanding (Reis et al., 2017).

9. A WEIGHT IMPLEMENTATION STANDARD

The word Weight should be used scientifically only when its use can be expanded into a declaration that makes the inference inspectable.

A proposed minimum standard is:

- 1. OBJECT — What already-defined thing is being said to have Weight?**
- 2. SYSTEM — For whom or what system?**
- 3. RELATION — Consequence for what measurable process, judgment, state, policy, or outcome?**
- 4. CONTEXT — Under what task, time, state, environment, comparison, or boundary?**
- 5. EVIDENCE — What observation or measure bears on the claimed relation?**
- 6. EVIDENTIAL STATUS — Is that evidence self-report, behavioral index, physiological measure, observer estimate, experimental manipulation, model inference, or another declared source?**
- 7. PROVENANCE — Where did the relevant organization come from, if source/history matters to the question?**
- 8. TRANSFER — If communication is involved, what changed in the receiver, and how was that change measured?**

Not every study needs all eight fields. Object, system, relation, context, and evidence are the minimum for a scientific Weight claim. Provenance and transfer are conditional fields that become necessary only when the inference depends on them. The claim level must also be explicit: descriptive claims require valid measurement of the declared relation, predictive claims

require demonstrated performance under a stated prediction/generalization regime, and causal claims require a valid causal-identification strategy.

When temporal route or processing level matters to the claim, the declaration must also state when the consequence was measured and at what observable level. A later self-report is evidence about the later reported state; it is not direct evidence about earlier processing without an independently justified bridge.

A complete declaration might look like this in ordinary prose:

"The warning was associated with stronger immediate avoidance policy in this participant during the driving task after direct exposure to the hazard, indexed by choice latency and route selection."

If the design supports causation, the language can be stronger:

"Experimentally increasing warning credibility increased avoidance-policy Weight, measured by route selection under the preregistered driving task."

Both are better than "the warning had high psychological Weight" because the object, system, relation, context, evidence, and inferential strength are visible.

Another example:

"Participants judged the speaker's statement as strongly grounded in direct experience, but this was an observer estimate of provenance and should not be treated as verified history."

And another:

"The diagnosis was accurately recalled but showed little relation to current treatment adherence during the first session."

That claim does not imply disbelief, lack of understanding, or lack of emotion unless those relations are separately measured.

The implementation standard succeeds only if it changes scientific practice. If expanding a Weight claim into these fields never changes the chosen measure, admissible inference, comparison, control, prediction, or failure condition, then the interface has not earned scientific status.

10. FORBIDDEN INFERENCES

The framework is designed to make several common shortcuts inadmissible.

Same words do not establish same Weight.

Same factual knowledge does not establish same understanding.

Same understanding does not establish same Weight.

A strong Weight estimate in one declared relation does not establish control of behavior.

High Weight does not establish truth.

Low Weight in a downstream relation does not establish disbelief.

Low or absent later report of phenomenal emotional consequence does not establish that a fully conscious emotion was present earlier and then ignored, dismissed, suppressed, or overridden.

Direct experience does not establish greater Weight.

Direct experience does not establish greater accuracy.

The same provenance does not establish the same present organization.

The same present consequence does not establish the same provenance.

Communication of content does not establish matched consequentiality in the receiver.

An observer estimate of Weight does not establish the target-side relation.

Perceived provenance does not establish verified provenance.

A measured association does not establish that X caused the relation attributed to Weight.

An action endpoint does not identify which represented objects participated, which carried relevant local consequence, or which lost competition.

A shared numerical score does not establish that unlike Weight measurements are commensurate.

These prohibitions are not philosophical decoration. Each changes what a researcher is allowed to conclude from a measurement.

11. MEASUREMENT WITHOUT A UNIVERSAL WEIGHT UNIT

An everyday word invites an everyday scale: no Weight, a little Weight, a lot of Weight. That can be useful in conversation, but it creates a scientific trap.

A score of 4 for a belief and a score of 4 for an emotion does not mean those objects contain equal amounts of a common psychological substance. Construct-validity work requires researchers to justify what a measure represents rather than infer identity from labels or numerical form (Flake & Fried, 2020; Lawson & Robins, 2021). Measurement harmonization makes the cross-construct point explicit: expressing scores on a common metric does not imply that the same construct is measured or that scores are substantively comparable across distinct constructs (de Beurs, Boehnke, & Fried, 2022).

The proposed rule is therefore:

**WEIGHT MAY BE MEASURED LOCALLY BEFORE IT IS COMPARED
GLOBALLY.**

Within a validated local relation, researchers should use the measurement technology appropriate to that domain: probability, reaction time, choice proportion, memory performance,

physiological response, willingness to pay, persistence, validated rating scales, or other measures.

The present paper does not propose a default 0-4, 0-5, interval, or ratio Weight scale. A generic ordinal scale would create apparent precision and invite cross-domain comparison before the necessary validation exists.

Researchers may use human-facing labels such as low, moderate, or high Weight when those labels are explicitly tied to locally defined thresholds or are clearly marked as informal description. Such labels do not create shared units.

Cross-object or cross-system magnitude comparison requires additional evidence. If a researcher wants to claim that a Belief.Weight score is quantitatively equivalent to an Emotion.Weight score, or that a sender's raw score equals a receiver's, a linking rule, measurement-invariance argument, calibration model, or other empirically justified bridge must be supplied. Equal display positions alone are insufficient.

The point of measurement is to force the question "Weight of what, for whom, in relation to what?" before inviting magnitude comparison.

12. MINIMAL FORMAL NOTATION

The paper does not require mathematics to define Weight. Formal notation is useful only after the human claim is clear.

When precision is needed, a Weight declaration can be written as:

$$W_R(X, S \mid \Omega)$$

where:

X is the already-typed object;

S is the target system;

R is the declared consequence relation;

Omega is the relevant context or boundary.

The notation does not imply that W is a scalar. Depending on the local science, the admissible description may be categorical, ordinal, continuous, multivariate, model-based, or simply a declared relation with an associated estimate and uncertainty.

This notation deliberately does not encode provenance. Provenance can be represented separately as $P(X, S)$ or as a structured source record because history and present consequence are not identical.

Likewise, an observer's estimate should be marked separately, for example:

$\hat{W}_{R,L}(X, S \mid \Omega)$

where L identifies the listener or observer making the estimate about target system S.

The notation does not add a mechanism. Its job is to stop an author from writing W when the object, system, relation, or context has not been declared.

If the notation does no work beyond that bookkeeping function, that is acceptable only if the bookkeeping changes an admissible inference, measurement, comparison, or test. If it changes nothing, the notation should be omitted.

13. IMPLEMENTATION IDEAS

13.1 Communication research

A communication study can independently measure what was said, what was understood, what the receiver believed, what changed in a specified consequence relation, and what the receiver inferred about the sender. This makes "successful communication" a profile rather than a single endpoint.

13.2 Risk and warning communication

Warnings are an obvious test domain because truth, comprehension, credibility, affect, and behavior can diverge. A warning can be readable but not credible, credible but low-stakes, high-stakes but poorly understood, or well understood yet overridden by competing goals. Weight declarations would force the study to specify which relation changed.

13.3 Clinical communication

A patient can accurately repeat a diagnosis while its practical consequences remain weakly integrated. Conversely, a diagnosis can carry enormous emotional Weight while the patient misunderstands its actual prognosis. Separating content, model, consequence, and source could improve description without assuming that one score captures all four.

13.4 Peer support and experiential knowledge

Peer support can be studied without claiming that shared experience automatically improves accuracy. Actual shared provenance, perceived shared provenance, felt understanding, regulation, trust, disclosure, and objective understanding can be measured separately.

13.5 Education

A learner may memorize a proposition without acquiring a generative model. A generative model may be acquired without the topic becoming consequential enough to guide later behavior. Weight can therefore be studied downstream of learning rather than used as another word for learning.

13.6 Expertise and professional judgment

Experts can acquire consequence through repeated model-based exposure rather than direct personal experience. This provides an important test of whether lived provenance is necessary for any claimed Weight relation. The framework predicts that it often will not be.

13.7 Human-AI interaction

A system may reproduce correct content without possessing the human provenance associated with that content, while a human listener may nevertheless assign Weight or authority to the output. The relevant scientific questions are then typed separately: content accuracy, model competence, source credibility, perceived understanding, and downstream influence.

14. EMPIRICAL TESTS

The proposal should survive only if its distinctions matter experimentally or methodologically. The first test therefore targets the interface itself rather than any one component phenomenon.

14.1 Test the Weight interface against generic explicitness

Select multiple domains in which the local constructs are already mature—for example warning communication, clinical communication, and education. For each domain, preregister the local constructs, evidence, comparison rules, and scoring criteria, then compare three analysis/reporting conditions:

- conventional reporting using the established domain constructs as normally practiced;
- a generic-explicitness control matched for effort and specificity, requiring construct identity, target system or population, measure, context, evidential status, inferential

strength, and comparison rule, but not organized through Weight or the content/model/consequence/provenance/transfer distinctions;

- the full Weight declaration standard.

The decisive comparison is Weight versus the matched generic-explicitness control, not Weight versus ordinary reporting. Outcomes should be scored prospectively for changes in measure selection, preregistered predictions, required controls, admissible comparisons, identification of observational equivalence, interpretation of null results, transfer claims, and explicit failure conditions.

The interface earns distinct scientific survival only if its organization exposes otherwise-missed dissociations, invalid comparisons, observational equivalences, transfer errors, or failure conditions beyond what generic methodological explicitness produces while remaining compatible with the local science. If the generic-explicitness control performs equivalently, Weight should contract to a human-language synthesis or design mnemonic rather than a distinct scientific interface, even if all underlying component phenomena are real.

14.2 Hold content constant and vary provenance

Teach or verify the same proposition through different routes: direct experience, vicarious experience, semantic instruction, professional modeling, or another controlled source. Then measure separately factual possession, generative understanding, and a predeclared local Weight relation.

The critical question is whether provenance predicts anything after present state has been adequately characterized. If not, provenance should remain source metadata rather than a required causal coordinate.

14.3 Cross verified provenance with perceived provenance

Manipulate whether two people actually share a relevant experience and whether the participant believes that they share it. Measure felt understanding, trust, disclosure, belonging, regulation, learning, and empathic accuracy.

If perceived shared history explains the social effects and verified history adds nothing, the mechanism belongs to shared-reality or perceived-similarity processes rather than to provenance itself.

14.4 Separate content, understanding, and Weight

Cross rote factual learning with generative explanatory learning and independently manipulate a known determinant of a predeclared local consequence relation. The design should make it possible to produce high factual knowledge with low understanding, high understanding with low local consequence, or high local consequence with poor understanding.

The manipulation should target the known local determinant rather than "Weight" as an unexplained latent cause. If the states cannot be dissociated under credible manipulations, the architecture should contract.

14.5 Test sender-receiver Weight mismatch directly

Measure the sender's declared local relation before communication. Measure the receiver's content acquisition, understanding, belief, and corresponding local relation after communication. The question is not whether the receiver "got the message" globally. It is which parts changed.

A strong demonstration would show accurate content transfer with little change in the receiver's corresponding relation, or a receiver-side consequence change without equivalent explanatory understanding.

Raw sender and receiver scores should not be treated as equal quantities unless the measure supports cross-person comparison. Within-person change, calibrated common measures, or explicitly linked scales are safer tests of mismatch.

14.6 Test the two forms of "it did not land"

Use propositions whose truth status can be independently established. Manipulate evidential support or source credibility separately from a determinant of downstream consequence such as current relevance, temporal distance, or task stakes. Measure belief confidence independently from the declared downstream relation.

This design distinguishes:

- failure to accept a proposition because the evidence or source carries insufficient Weight for belief revision;
- acceptance of the proposition with weak consequence in another declared relation.

Evidence that the two manipulations selectively affect different outcomes would support the separation. If belief confidence alone accounts for the downstream pattern across credible domains, the broader uptake distinction should contract.

14.7 Test observer estimates

Use targets with independently verified current states and, where possible, independently verified provenance classes. Present transcript-only, audio-only, video, and multimodal conditions. Test whether observers can estimate current consequence or provenance above

chance while controlling for confidence, charisma, expressivity, arousal, similarity, personality, and social skill.

Failure to infer provenance would not kill Weight. It would kill only the stronger claim that provenance is naturally readable from behavior.

15. RELATION TO EXISTING SCIENCE

The proposed framework intentionally uses established science rather than competing with it.

Communication theory already separates several relations that ordinary language compresses. Speech-act theory distinguishes propositional content from what an utterance does and the effects it produces (Austin, 1962). Grounding and common-ground work treats communication as coordinated evidence of mutual understanding rather than identity of internal states (Clark & Brennan, 1991). McGuire's communication-persuasion matrix separates exposure, attention, comprehension, acceptance, retention, and action (McGuire, 1981). These traditions strongly establish prior precedent for the claim that receiving words is not the same as every downstream effect of communication.

Relevance theory is a particularly close communication-level precedent because it distinguishes linguistically decoded content from inferential processing and defines relevance for an individual partly by the cognitive effects an input yields relative to processing effort (Sperber & Wilson, 1987). The persuasion literature on personal involvement likewise shows that the same message can be processed differently as a function of receiver-relative relevance, with high involvement shifting sensitivity toward argument quality (Petty, Cacioppo, & Goldman, 1981). These frameworks strongly precedent the idea that information can be available while differing in recipient-local consequence. They do not, by themselves, supply the full cross-domain

identification architecture proposed here: typed parent objects, content versus generative understanding, present consequence versus provenance, target-side evidence versus observer estimate, noncommensurate local measurement, and typed transfer.

Risk communication provides another set of local precedents. The social-amplification-of-risk framework models how risk signals can be amplified or attenuated through transfer and response systems (Kasperson et al., 1988). Risk-as-feelings work separates cognitive assessment from affective response and shows that the two can diverge (Loewenstein et al., 2001). These literatures already own important examples of information and consequence coming apart.

Attitude-strength research shows that apparently simple "strength" language decomposes into multiple measurable properties with different antecedents and consequences (Krosnick et al., 1993; Howe & Krosnick, 2017).

Regulatory engagement theory shows that experienced value intensity can depend on strength of engagement in addition to hedonic direction (Higgins, 2006). This is a close precedent for the idea that what something means and how forcefully it matters should not be treated as one thing.

Shared-reality theory shows that perceived commonality in inner states about a target can have relational and epistemic functions (Echterhoff et al., 2009). Research on felt understanding further distinguishes the experience of being understood from objective or externally assessed understanding (Reis et al., 2017). These findings establish prior precedent for important parts of the interpersonal side of Weight.

Epistemic-vigilance theory explains why communicated propositions are evaluated rather than automatically accepted (Sperber et al., 2010), providing established machinery for one route by which a statement can fail to be accepted as immediate fact.

Source-monitoring research establishes that origin/source is separable from content and that source judgments are fallible (Johnson et al., 1993).

Research on explanatory depth separates possession of facts and familiarity from the ability to produce mechanistic explanation (Rozenblit & Keil, 2002).

Empathy research shows that shared experience, perceived understanding, concern, and actual accuracy need not move together (Hodges et al., 2010; Israelashvili et al., 2020; Hall & Schmid Mast, 2007).

Measurement research requires explicit construct validation and warns against conflating related but distinct constructs or treating a common metric as proof of construct equivalence (Flake & Fried, 2020; Lawson & Robins, 2021; de Beurs et al., 2022).

A targeted same-level prior-art recheck therefore treats relevance theory, personal-involvement and persuasion models, attitude-strength and regulatory-engagement traditions, shared reality, epistemic vigilance, source monitoring, and construct-validity frameworks as serious precedents rather than as background decoration. No one of these frameworks, as reviewed here, was shown to assemble substantially the same full identification function. That is not evidence that paper-level novelty is confirmed; it leaves the paper-level novelty question OPEN and testable against the interface comparison in Section 14.1.

These precedents do not need to be novel for Weight to use them. Their role is to defend the human-language distinctions and constrain what Weight is allowed to mean.

The paper-level question is therefore not whether psychology, communication science, decision science, or social psychology has previously studied value, importance, understanding, persuasion, empathy, source, credibility, or transfer. It plainly has. The question is whether a single human-readable, typed identification standard is useful for keeping those established

relations from being silently collapsed when people say that something "has Weight" or expect another person to "understand its Weight."

That larger contribution remains a research proposition rather than an established fact. It survives only if the implementation standard changes scientific work in the ways specified above.

16. LIMITS, CONTRACTION, AND KILL CONDITIONS

Weight should not survive merely because the word is intuitive.

If every scientifically admissible Weight claim is fully captured by an established local construct and the interface changes no inference, comparison, measurement, control, prediction, or communication standard, Weight should contract to informal language.

If one existing same-level framework already provides the full typed architecture described here—object identity, local consequence, provenance, target versus observer estimate, content versus generative understanding, and typed communication outcomes—then the paper-level contribution is subsumed even though the human examples remain valid.

If the direct interface test shows that adding the Weight declaration standard produces no reliable difference in measurement design, inferential discipline, comparison, or detection of dissociation across credible domains, then the proposed scientific interface has failed.

If provenance adds no predictive or intervention value after present state is adequately characterized, provenance should not be treated as an independent causal coordinate.

If sender-receiver Weight mismatch reduces entirely to already-measured differences in comprehension, belief confidence, or a single local construct in every tested domain, the general transfer distinction should contract accordingly.

If the two forms of "it did not land" cannot be separated—if belief confidence fully accounts for the downstream consequence patterns across credible domains—then the immediate-fact argument should contract to epistemic acceptance rather than retain a separate uptake claim.

If observer estimates of provenance fail, the framework should retain the distinction between target-side and perceived Weight while abandoning claims that provenance can be reliably read from natural behavior.

Failure of a universal Weight scale does not kill the typed framework because no universal scale is claimed. It forbids universal commensurability. Local measures should remain local unless a bridge is empirically justified.

Most importantly, the paper must never rescue Weight by making it unfalsifiable. "Weight" cannot mean whatever remains unexplained after measurement. Every scientific use must commit to an object, system, relation, context, evidential status, and appropriate failure condition before the outcome is interpreted.

17. CONCLUSION

People already know the phenomenon this paper is trying to describe.

They know what it is like to say something important and realize the other person heard the sentence but did not receive what it carried.

They know what it is like to be told something true and only later discover that the fact had not yet reorganized the world they were using to think, feel, or act.

They know that a person who has lived through something may speak from a different place than a person who has only studied it, while also knowing that lived experience does not automatically make someone right.

They know that "I understand" can mean I recognize the words, I can explain the model, I have felt something like this, I see why it matters to you, it matters to me too, or simply I am trying to tell you that I am listening.

Science should not remove those distinctions merely because ordinary language compresses them. Nor should it turn the compression into a new universal mechanism.

Weight is proposed as a disciplined middle layer.

Something has Weight only as an already-defined object, for a particular system, in relation to a particular consequence, under particular conditions. "Consequence" does not itself establish causation. Its history may matter, but history is not present consequence. Its truth may matter, but truth is not Weight. Its Weight may bear on behavior, but Weight is not behavior. Another person may estimate it, but an observer estimate is not the target-side relation. It may change during communication, but content does not carry consequence automatically from one system to another.

The framework therefore makes two practical claims. First, a sender cannot infer matched Weight from successful transmission of content. Second, a receiver's failure to treat a statement as immediate reality can reflect either insufficient evidential support for acceptance or weak consequential uptake after acceptance; those cases must be measured separately.

The simplest version of the proposal remains the most important:

Having the words, having the model, having the consequence, and transmitting any of them are different things.

A scientifically disciplined use of Weight gives those differences a common language without pretending they are the same mechanism. Whether that common language deserves

scientific survival now has a direct test: it must change what researchers measure, compare, infer, or predict. If it does not, Weight should remain only the useful human word it already is.

References

- Austin, J. L. (1962). *How to do things with words* (J. O. Urmson, Ed.). Oxford University Press.
- Berridge, K. C., Robinson, T. E., & Aldridge, J. W. (2009). Dissecting components of reward: “liking”, “wanting”, and learning. *Current Opinion in Pharmacology*, 9(1), 65–73.
<https://doi.org/10.1016/j.coph.2008.12.014>
- Boden, M. T., & Thompson, R. J. (2017). Meta-analysis of the association between emotional clarity and attention to emotions. *Emotion Review*, 9(1), 79–85.
<https://doi.org/10.1177/1754073915610640>
- Clark, H. H., & Brennan, S. E. (1991). Grounding in communication. In L. B. Resnick, J. M. Levine, & S. D. Teasley (Eds.), *Perspectives on socially shared cognition* (pp. 127–149). American Psychological Association. <https://doi.org/10.1037/10096-006>
- de Beurs, E., Boehnke, J. R., & Fried, E. I. (2022). Common measures or common metrics? A plea to harmonize measurement results. *Clinical Psychology & Psychotherapy*, 29(5), 1755–1767. <https://doi.org/10.1002/cpp.2742>
- Echterhoff, G., Higgins, E. T., & Levine, J. M. (2009). Shared reality: Experiencing commonality with others’ inner states about the world. *Perspectives on Psychological Science*, 4(5), 496–521. <https://doi.org/10.1111/j.1745-6924.2009.01161.x>
- Flake, J. K., & Fried, E. I. (2020). Measurement schmeasurement: Questionable measurement practices and how to avoid them. *Advances in Methods and Practices in Psychological Science*, 3(4), 456–465. <https://doi.org/10.1177/2515245920952393>

- Gross, J. J. (1998). Antecedent- and response-focused emotion regulation: Divergent consequences for experience, expression, and physiology. *Journal of Personality and Social Psychology*, 74(1), 224–237. <https://doi.org/10.1037/0022-3514.74.1.224>
- Gyurak, A., Gross, J. J., & Etkin, A. (2011). Explicit and implicit emotion regulation: A dual-process framework. *Cognition and Emotion*, 25(3), 400–412. <https://doi.org/10.1080/02699931.2010.544160>
- Hall, J. A., & Schmid Mast, M. (2007). Sources of accuracy in the empathic accuracy paradigm. *Emotion*, 7(2), 438–446. <https://doi.org/10.1037/1528-3542.7.2.438>
- Higgins, E. T. (2006). Value from hedonic experience and engagement. *Psychological Review*, 113(3), 439–460. <https://doi.org/10.1037/0033-295X.113.3.439>
- Hodges, S. D., Kiel, K. J., Kramer, A. D. I., Veach, D., & Villanueva, B. R. (2010). Giving birth to empathy: The effects of similar experience on empathic accuracy, empathic concern, and perceived empathy. *Personality and Social Psychology Bulletin*, 36(3), 398–409. <https://doi.org/10.1177/0146167209350326>
- Howe, L. C., & Krosnick, J. A. (2017). Attitude strength. *Annual Review of Psychology*, 68, 327–351. <https://doi.org/10.1146/annurev-psych-122414-033600>
- Israelashvili, J., Sauter, D. A., & Fischer, A. H. (2020). Different faces of empathy: Feelings of similarity disrupt recognition of negative emotions. *Journal of Experimental Social Psychology*, 87, 103912. <https://doi.org/10.1016/j.jesp.2019.103912>
- Johnson, M. K., Hashtroudi, S., & Lindsay, D. S. (1993). Source monitoring. *Psychological Bulletin*, 114(1), 3–28. <https://doi.org/10.1037/0033-2909.114.1.3>

- Kasperson, R. E., Renn, O., Slovic, P., Brown, H. S., Emel, J., Goble, R., Kasperson, J. X., & Ratick, S. (1988). The social amplification of risk: A conceptual framework. *Risk Analysis*, 8(2), 177–187. <https://doi.org/10.1111/j.1539-6924.1988.tb01168.x>
- Krosnick, J. A., Boninger, D. S., Chuang, Y. C., Berent, M. K., & Carnot, C. G. (1993). Attitude strength: One construct or many related constructs? *Journal of Personality and Social Psychology*, 65(6), 1132–1151. <https://doi.org/10.1037/0022-3514.65.6.1132>
- Lawson, K. M., & Robins, R. W. (2021). Sibling constructs: What are they, why do they matter, and how should you handle them? *Personality and Social Psychology Review*, 25(4), 344–366. <https://doi.org/10.1177/10888683211047101>
- Loewenstein, G. F., Weber, E. U., Hsee, C. K., & Welch, N. (2001). Risk as feelings. *Psychological Bulletin*, 127(2), 267–286. <https://doi.org/10.1037/0033-2909.127.2.267>
- McGuire, W. J. (1981). Theoretical foundations of campaigns. In R. E. Rice & W. J. Paisley (Eds.), *Public communication campaigns* (pp. 41–70). Sage.
- Petty, R. E., Cacioppo, J. T., & Goldman, R. (1981). Personal involvement as a determinant of argument-based persuasion. *Journal of Personality and Social Psychology*, 41(5), 847–855. <https://doi.org/10.1037/0022-3514.41.5.847>
- Reis, H. T., Lemay, E. P., Jr., & Finkenauer, C. (2017). Toward understanding understanding: The importance of feeling understood in relationships. *Social and Personality Psychology Compass*, 11(3), e12308. <https://doi.org/10.1111/spc3.12308>
- Rozenblit, L., & Keil, F. (2002). The misunderstood limits of folk science: An illusion of explanatory depth. *Cognitive Science*, 26(5), 521–562. https://doi.org/10.1207/S15516709COG2605_1

Sperber, D., & Wilson, D. (1987). *Précis of Relevance: Communication and cognition*.

Behavioral and Brain Sciences, 10(4), 697–710.

<https://doi.org/10.1017/S0140525X00055345>

Sperber, D., Clément, F., Heintz, C., Mascaro, O., Mercier, H., Origgi, G., & Wilson, D. (2010).

Epistemic vigilance. *Mind & Language*, 25(4), 359–393. <https://doi.org/10.1111/j.1468->

0017.2010.01394.x